

```
In [ ]: # Importing Libraries
```

```
import numpy as np
import matplotlib.pyplot as plt
import keras
from keras.datasets import mnist

from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout
```

```
In [ ]: mnist.load_data?
```

```
In [ ]: (X_train, y_train), (X_test, y_test) = mnist.load_data()
```


```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/
mnist.npz
11490434/11490434 ━━━━━━━━ 1s 0us/step
```



```
In [ ]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```



```
Out[ ]: ((60000, 28, 28), (60000,), (10000, 28, 28), (10000,))
```



```
In [ ]: def plot_input_img(i):
    plt.imshow(X_train[i], cmap = 'binary')
    plt.title(y_train[i])
    plt.show()
```



```
In [ ]: for i in range(10):
    plot_input_img(i)
```














```
In [ ]: from os import XATTR_SIZE_MAX
# Preprocessing the images

# Normalizing the images to [0,1] scale
X_train = X_train.astype(np.float32)/255
X_test = X_test.astype(np.float32)/255

# Reshape or expand the dimensions of the images to (28,28,1)
X_train = np.expand_dims(X_train, -1)
X_test = np.expand_dims(X_test, -1)

# Converting classes to one hot vectors
y_train = keras.utils.to_categorical(y_train)
y_test = keras.utils.to_categorical(y_test)
```

```
In [ ]: X_train.shape
```

```
Out[ ]: (60000, 28, 28, 1, 1)
```

```
In [ ]: # Model building

model = Sequential()

model.add(Conv2D(32, (3,3), input_shape = (28,28,1), activation= 'relu'))
model.add(MaxPool2D((2,2)))

model.add(Conv2D(64, (3,3), activation= 'relu'))
model.add(MaxPool2D((2,2)))

model.add(Flatten())
model.add(Dropout(0.25))

model.add(Dense(10, activation= "softmax"))
```

```
/usr/local/lib/python3.12/dist-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer.
When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

In []: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dropout (Dropout)	(None, 1600)	0
dense (Dense)	(None, 10)	16,010

Total params: 34,826 (136.04 KB)

Trainable params: 34,826 (136.04 KB)

Non-trainable params: 0 (0.00 B)

In []: `model.compile(optimizer= 'adam', loss = keras.losses.categorical_crossentropy, metr`

In []: `# Callbacks`

```
from keras.callbacks import EarlyStopping, ModelCheckpoint

# EarlyStopping

es = EarlyStopping(monitor='val_acc', min_delta=0.01, patience= 4, verbose= 1)

# Model check point

mc = ModelCheckpoint("./bestmodel.h5", monitor= "val_acc", verbose= 1, save_best_only=True)

cb = [es,mc]
```

In []: `# Model training`

```
his = model.fit(X_train, y_train, epochs = 5, validation_split= 0.3)
```

```
Epoch 1/5
1313/1313 38s 29ms/step - accuracy: 0.9480 - loss: 0.1698 - val_accuracy: 0.9632 - val_loss: 0.1235
Epoch 2/5
1313/1313 38s 29ms/step - accuracy: 0.9529 - loss: 0.1524 - val_accuracy: 0.9663 - val_loss: 0.1154
Epoch 3/5
1313/1313 38s 29ms/step - accuracy: 0.9586 - loss: 0.1420 - val_accuracy: 0.9682 - val_loss: 0.1082
Epoch 4/5
1313/1313 40s 29ms/step - accuracy: 0.9590 - loss: 0.1381 - val_accuracy: 0.9686 - val_loss: 0.1069
Epoch 5/5
1313/1313 38s 29ms/step - accuracy: 0.9599 - loss: 0.1334 - val_accuracy: 0.9711 - val_loss: 0.0993
```

```
In [ ]: model_S = keras.models.load_model('my_bestmodel.h5')
```

```
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.
```

```
In [ ]: score = model_S.evaluate(X_test, y_test)
```

```
print(f"My model accuracy is {score[1]}")
```

```
313/313 3s 8ms/step - accuracy: 0.9700 - loss: 0.0975
My model accuracy is 0.9750000238418579
```

```
In [ ]:
```